

# Q Learning and Deep Q Network

Ling Fei Zhang  
260985358


April 17, 2023

# Table of Contents


- 1 Introduction
- 2 Methodology
- 3 Results
- 4 Conclusion

# Games

CartPole



Lunar Lander



# Agents

- Q Learning
  - Model-free reinforcement learning algorithm
  - Uses a table to store Q values for each state-action pair
  - Effective in simple environments
  - Struggles in more complex environments since it is impractical to store and update Q values for all the state-action pairs
- Deep Q Network (DQN)
  - Uses neural networks to learn policies to map states to Q values
  - Neural networks can handle large state spaces and continuous action spaces

# Agents

- Q Learning
  - Model-free reinforcement learning algorithm
  - Uses a table to store Q values for each state-action pair
  - Effective in simple environments
  - Struggles in more complex environments since it is impractical to store and update Q values for all the state-action pairs
- Deep Q Network (DQN)
  - Uses neural networks to learn policies to map states to Q values
  - Neural networks can handle large state spaces and continuous action spaces

# Agents

- Q Learning
  - Model-free reinforcement learning algorithm
  - Uses a table to store Q values for each state-action pair
  - Effective in simple environments
  - Struggles in more complex environments since it is impractical to store and update Q values for all the state-action pairs
- Deep Q Network (DQN)
  - Uses neural networks to learn policies to map states to Q values
  - Neural networks can handle large state spaces and continuous action spaces

# Q Learning

- Impractical to store all Q values for a continuous observation space
- Discretization needed
- Hyperparameters used for Q Learning
  - Discount factor  $\gamma = 0.99$
  - Start epsilon of 1
  - End epsilon of 0.001
  - Epsilon decay of  $5^{-4}$
  - 10 bins
  - learning rate of 0.25 and 0.005 for CartPole and Lunar Lander, respectively

# Q Learning

- Impractical to store all Q values for a continuous observation space
- Discretization needed
- Hyperparameters used for Q Learning
  - Discount factor  $\gamma = 0.99$
  - Start epsilon of 1
  - End epsilon of 0.001
  - Epsilon decay of  $5^{-4}$
  - 10 bins
  - learning rate of 0.25 and 0.005 for CartPole and Lunar Lander, respectively



# Q Learning

- Impractical to store all Q values for a continuous observation space
- Discretization needed
- Hyperparameters used for Q Learning
  - Discount factor  $\gamma = 0.99$
  - Start epsilon of 1
  - End epsilon of 0.001
  - Epsilon decay of  $5^{-4}$
  - 10 bins
  - learning rate of 0.25 and 0.005 for CartPole and Lunar Lander, respectively

# Q Learning

- Impractical to store all Q values for a continuous observation space
- Discretization needed
- Hyperparameters used for Q Learning
  - Discount factor  $\gamma = 0.99$
  - Start epsilon of 1
  - End epsilon of 0.001
  - Epsilon decay of  $5^{-4}$
  - 10 bins
  - learning rate of 0.25 and 0.005 for CartPole and Lunar Lander, respectively

# Q Learning Algorithm

---

## Algorithm Q Learning(episodes, $\alpha, \epsilon, \gamma$ )

---

- 1: Initialize  $Q(s, a)$  for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$  arbitrarily
  - 2: Set  $Q(\text{terminal}, \cdot) = 0$  for all terminal states
  - 3: **for each** episode in episodes **do**
  - 4:     Initialize  $s$
  - 5:      $done \leftarrow \text{False}$
  - 6:     **while** not  $done$  **do**
  - 7:         Choose  $a \in \mathcal{A}$  from  $s$  using policy derived from  $Q$
  - 8:         Take action  $a$  and observe reward  $r$  and next state  $s'$
  - 9:          $Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_a Q(s', a) - Q(s, a)]$
  - 10:         $s \leftarrow s'$
  - 11:     **end while**
  - 12: **end for**
-

# DQN Algorithm Setup

- Implementation of ReplayMemory, which acts as a replay buffer
- Implementation of deep neural network
  - Two hidden layers, each with 128 units
  - ReLU activation function
  - Adam's optimizer
  - Huber loss

# DQN Algorithm Setup

- Implementation of ReplayMemory, which acts as a replay buffer
- Implementation of deep neural network
  - Two hidden layers, each with 128 units
  - ReLU activation function
  - Adam's optimizer
  - Huber loss

# DQN Algorithm

---

**Algorithm** DQN(episodes,  $\alpha, \epsilon, \gamma, C$ )

---

```
1: Initialize replay buffer  $\mathcal{D}$  with maximum capacity 100000
2: Initialize policy network  $Q$  with random weights  $\theta$ 
3: Initialize target network  $\tilde{Q}$  with random weights  $\tilde{\theta}$ 
4: for each episode in episodes do
5:   Initialize  $s$ 
6:   Set  $t \leftarrow 0$ 
7:   done  $\leftarrow$  False
8:   while not done do
9:     Choose  $a \in \mathcal{A}$  from  $s$  using policy network  $Q$ 
10:    Take action  $a$  and observe reward  $r$ , next state  $s'$  and done
11:    Store transition  $(s, a, r, s', \text{done})$ 
12:    Sample a minibatch of random transitions  $(s, a, r, s', \text{done})$  from  $\mathcal{D}$ 
13:    Set  $\hat{y} = \begin{cases} r & \text{if } s \text{ is a terminal state} \\ r + \gamma \max_a \tilde{Q}(s', a; \tilde{\theta}) & \text{otherwise} \end{cases}$ 
14:    Perform gradient descent on  $(\hat{y} - Q(s, a; \theta))^2$  w.r.t. the policy
    network parameters  $\theta$ 
15:    if  $t \bmod C = 0$  then
16:       $\tilde{Q} \leftarrow Q$ 
17:    end if
18:  end while
19: end for
```

---

- Discount factor  $\gamma = 0.99$
- Learning rate 0.003
- Start epsilon of 1
- End epsilon of 0.01
- Epsilon decay of  $5^{-4}$
- $\tau = 0.005$
- Default batch size of 64

# DQN Algorithm

---


**Algorithm** DQN(episodes,  $\alpha, \epsilon, \gamma, C$ )

---


- 1: Initialize replay buffer  $\mathcal{D}$  with maximum capacity 100000
  - 2: Initialize policy network  $Q$  with random weights  $\theta$
  - 3: Initialize target network  $\tilde{Q}$  with random weights  $\tilde{\theta}$
  - 4: **for** each episode in episodes **do**
  - 5:     Initialize  $s$
  - 6:     Set  $t \leftarrow 0$
  - 7:      $done \leftarrow \text{False}$
  - 8:     **while** not  $done$  **do**
  - 9:         Choose  $a \in \mathcal{A}$  from  $s$  using policy network  $Q$
  - 10:        Take action  $a$  and observe reward  $r$ , next state  $s'$  and  $done$
  - 11:        Store transition  $(s, a, r, s', done)$
  - 12:        Sample a minibatch of random transitions  $(s, a, r, s', done)$  from  $\mathcal{D}$
  - 13:        Set  $\hat{y} = \begin{cases} r & \text{if } s \text{ is a terminal state} \\ r + \gamma \max_a \tilde{Q}(s', a; \tilde{\theta}) & \text{otherwise} \end{cases}$
  - 14:        Perform gradient descent on  $(\hat{y} - Q(s, a; \theta))^2$  w.r.t. the policy network parameters  $\theta$
  - 15:        **if**  $t \bmod C = 0$  **then**
  - 16:           $\tilde{Q} \leftarrow Q$
  - 17:        **end if**
  - 18:     **end while**
  - 19: **end for**
- 

- Discount factor  $\gamma = 0.99$
- Learning rate 0.003
- Start epsilon of 1
- End epsilon of 0.01
- Epsilon decay of  $5^{-4}$
- $\tau = 0.005$
- Default batch size of 64

# Plots



(a) CartPole



(b) Lunar Lander

**Figure:** 500 training episodes for Q Learning and DQN. Returns are averaged over the last 100 episodes.




# CartPole

## DQN

- Learns much quicker than Q Learning
- Reached an average return of 300 at around 300 episodes
- Unstable learning process: significant drop at episode 350
- Agent is possibly trying to escape a local minimum

## Q Learning

- Learns much slower compared to DQN
- Not great returns, even after 500 training episodes
- Agent might be stuck in a local minima




# CartPole

## DQN

- Learns much quicker than Q Learning
- Reached an average return of 300 at around 300 episodes
- Unstable learning process: significant drop at episode 350
- Agent is possibly trying to escape a local minimum

## Q Learning

- Learns much slower compared to DQN
- Not great returns, even after 500 training episodes
- Agent might be stuck in a local minima




# CartPole

## DQN

- Learns much quicker than Q Learning
- Reached an average return of 300 at around 300 episodes
- Unstable learning process: significant drop at episode 350
- Agent is possibly trying to escape a local minimum

## Q Learning

- Learns much slower compared to DQN
- Not great returns, even after 500 training episodes
- Agent might be stuck in a local minima




# Lunar Lander

## DQN

- DQN outperforms Q Learning
- Unstable learning process

## Q Learning

- Slow but stable learning
- No drastic change in average return.  
Instead, the average return slowly increases




# Lunar Lander

## DQN

- DQN outperforms Q Learning
- Unstable learning process

## Q Learning

- Slow but stable learning
- No drastic change in average return.  
Instead, the average return slowly increases




# Lunar Lander

## DQN

- DQN outperforms Q Learning
- Unstable learning process

## Q Learning

- Slow but stable learning
- No drastic change in average return.  
Instead, the average return slowly increases



# Conclusion

## DQN

- Learns much faster in both games, but the agent exhibited some instability during the learning process

## Q Learning

- Learns much slower, but more steadily

The choice of algorithm depends on the specific application and trade-offs between speed and stability.

# Conclusion

## DQN

- Learns much faster in both games, but the agent exhibited some instability during the learning process

## Q Learning

- Learns much slower, but more steadily

The choice of algorithm depends on the specific application and trade-offs between speed and stability.



# Conclusion

## DQN

- Learns much faster in both games, but the agent exhibited some instability during the learning process

## Q Learning

- Learns much slower, but more steadily

The choice of algorithm depends on the specific application and trade-offs between speed and stability.

# Conclusion

## DQN

- Learns much faster in both games, but the agent exhibited some instability during the learning process

## Q Learning

- Learns much slower, but more steadily

The choice of algorithm depends on the specific application and trade-offs between speed and stability.

# Conclusion

## DQN

- Learns much faster in both games, but the agent exhibited some instability during the learning process

## Q Learning

- Learns much slower, but more steadily

**The choice of algorithm depends on the specific application and trade-offs between speed and stability.**

# Future Work

## Deep deterministic policy gradient (DDPG)

- Combines Q Learning with policy gradients to learn a deterministic policy directly
- This has shown to be effective in continuous action spaces and could address some of the instability issues observed in DQN.

# Future Work

## Deep deterministic policy gradient (DDPG)

- Combines Q Learning with policy gradients to learn a deterministic policy directly
- This has shown to be effective in continuous action spaces and could address some of the instability issues observed in DQN.

# Future Work

## Deep deterministic policy gradient (DDPG)

- Combines Q Learning with policy gradients to learn a deterministic policy directly
- This has shown to be effective in continuous action spaces and could address some of the instability issues observed in DQN.